



M17 Protocol Specification

DRAFT

M17 Working Group:

Wojciech	SP5WWP
Juhani	OH1CAU
Elms	KM6VMZ
Nikoloz	SO3ALG
Mark	KR6ZY
Steve	KC1AWV
Rob	WX9O
Tom	N7TAE
Mike	W2FBI

Jul 04, 2021

Published July 4, 2021

Copyright © 2019-2021 M17 Working Group

Permission is granted to make and distribute verbatim copies of this document provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this document under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this document into another language, under the above conditions for modified versions, except that this permission notice may be included in translations approved by the Free Software Foundation instead of in the original English.

See the GNU General Public License version 2 for more details.

Part I - Air and IP Interface:

1	M17 RF Protocol: Summary	1
2	Glossary	3
3	Physical Layer	5
3.1	4FSK generation	5
3.2	Preamble	6
3.3	Bit types	6
3.4	Error correction coding schemes and bit type conversion	6
4	Data Link Layer	11
4.1	Stream Mode	11
4.2	Packet Mode	15
5	Application Layer	19
5.1	Packet Superframes	19
5.2	Encryption Types	20
A	Address Encoding	23
A.1	Callsign Encoding: base40	23
A.2	Callsign Formats	25
B	Decorrelator sequence	27
C	Interleaving	29
D	M17 Internet Protocol (IP) Networking	33
D.1	Standard IP Framing	33
D.2	Control Packets	34
E	KISS Protocol	37
E.1	References	37
E.2	Glossary	37
E.3	M17 Protocols	38
E.4	KISS Basics	38
E.5	Packet Protocols	39
E.6	Stream Protocol	40
E.7	Mixing Modes	42

E.8 Implementation Details	42
Bibliography	43
Index	45

M17 RF Protocol: Summary

M17 is an RF protocol that is:

- Completely open: open specification, open source code, open source hardware, open algorithms. Anyone must be able to build an M17 radio and interoperate with other M17 radios without having to pay anyone else for the right to do so.
- Optimized for amateur radio use.
- Simple to understand and implement.
- Capable of doing the things hams expect their digital protocols to do:
 - Voice (eg: DMR, D-Star, etc)
 - Point to point data (eg: Packet, D-Star, etc)
 - Broadcast telemetry (eg: APRS, etc)
 - Extensible, so more capabilities can be added over time.

To do this, the M17 protocol is broken down into three protocol layers, like a network:

1. Physical Layer: How to encode 1s and 0s into RF. Specifies RF modulation, symbol rates, bits per symbol, etc.
2. Data Link Layer: How to packetize those 1s and 0s into usable data. Packet vs Stream modes, headers, addressing, etc.
3. Application Layer: Accomplishing activities. Voice and data streams, control packets, beacons, etc.

This document attempts to document these layers.

ECC Error Correcting Code

FEC Forward Error Correction

Frame The individual components of a stream, each of which contains payload data interleaved with frame signalling.

Link Information Frame The first frame of any transmission. It contains full LICH data.

LICH Link Information Channel. The LICH contains all information needed to establish an M17 link. The first frame of a transmission contains full LICH data, and subsequent frames each contain one sixth of the LICH data so that late-joiners can obtain the LICH.

Packet A single burst of transmitted data containing 100s to 1000s of bytes, after which the physical layer stops sending data.

Superframe A set of six consecutive frames which collectively contain full LICH data are grouped into a superframe.

3.1 4FSK generation

M17 standard uses 4FSK modulation running at 4800 symbols/s (9600 bits/s) with a deviation index $h=0.33$ for transmission in 9 kHz channel bandwidth. Channel spacing is 12.5 kHz. The symbol stream is converted to a series of impulses which pass through a root-raised-cosine ($\alpha=0.5$) shaping filter before frequency modulation at the transmitter and again after frequency demodulation at the receiver.



Fig. 3.1: 4FSK modulator dataflow

The bit-to-symbol mapping is shown in the table below.

Table 3.1: Dibit symbol mapping to 4FSK deviation

Information bits		Symbol	4FSK deviation
Bit 1	Bit 0		
0	1	+3	+2.4 kHz
0	0	+1	+0.8 kHz
1	0	-1	-0.8 kHz
1	1	-3	-2.4 kHz

The most significant bits are sent first, meaning that the byte 0xB4 (= 0b10'11'01'00) in type 4 bits (see *Bit types*) would be sent as the symbols -1 -3 +3 +1.

3.2 Preamble

Every transmission starts with a preamble, which shall consist of at least 40 ms of alternating -3, +3... symbols. This is equivalent to 40 milliseconds of a 2400 Hz tone

3.3 Bit types

The bits at different stages of the error correction coding are referred to with bit types, given in [Table 3.2](#).

Table 3.2: Bit types

Type 1	Data link layer bits
Type 2	Bits after appropriate encoding
Type 3	Bits after puncturing (only for convolutionally coded data, for other ECC schemes type 3 bits are the same as type 2 bits)
Type 4	Decorrelated and interleaved (re-ordered) type 3 bits

Type 4 bits are used for transmission over the RF. Incoming type 4 bits shall be decoded to type 1 bits, which are then used to extract all the frame fields.

3.4 Error correction coding schemes and bit type conversion

Two distinct *ECC/FEC* schemes are used for different parts of the transmission.

3.4.1 Link setup frame (LSF)

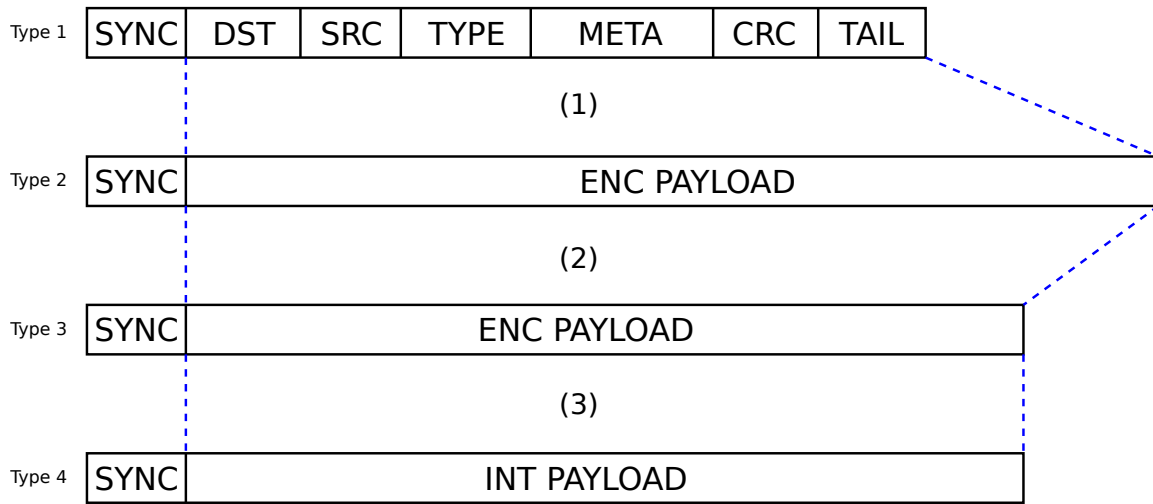


Fig. 3.2: ECC stages for the link setup frame

240 DST, SRC, TYPE, META and CRC type 1 bits are convolutionally coded using rate 1/2 coder with constraint $K=5$. 4 tail bits are used to flush the encoder's state register, giving a total of 244 bits being encoded. Resulting 488 type 2 bits are retained for type 3 bits computation. Type 3 bits are computed by puncturing type 2 bits using a scheme shown in chapter 4.4. This results in 368 bits, which in conjunction with the synchronization burst gives 384 bits ($384 \text{ bits} / 9600\text{bps} = 40 \text{ ms}$).

Interleaving type 3 bits produce type 4 bits that are ready to be transmitted. Interleaving is used to combat error bursts.

3.4.2 Subsequent frames

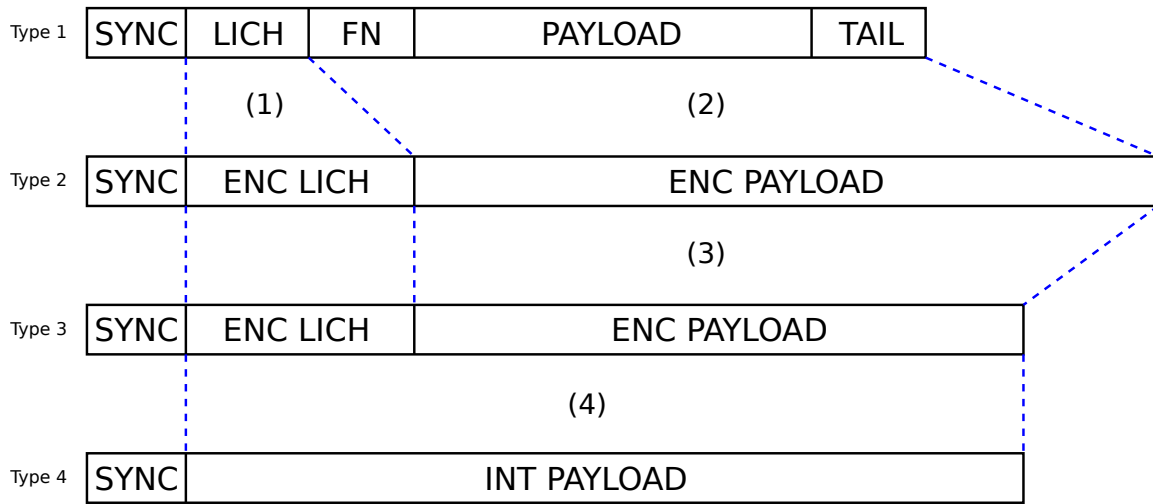


Fig. 3.3: ECC stages of subsequent frames

A 48-bit (type 1) chunk of the LSF is partitioned into 4 12-bit parts and encoded using Golay (24, 12) code. This produces 96 encoded bits of type 2. These bits are used in the Link Information Channel (LICH).

16-bit FN and 128 bits of payload (144 bits total) are convolutionally encoded in a manner analogous to that of the link setup frame. A total of 148 bits is being encoded resulting in 296 type 2 bits. These bits are punctured to generate 272 type 3 bits.

96 type 2 bits of LICH are concatenated with 272 type 3 bits and re-ordered to form type 4 bits for transmission. This, along with 16-bit sync in the beginning of frame, gives a total of 384 bits

The LICH chunks allow for late listening and independent decoding to check destination address. The goal is to require less complexity to decode just the LICH and check if the full message should be decoded.

3.4.3 Golay (24,12)

The Golay (24,12) encoder uses the polynomial 0xC75 to generate the 11 check bits. The check bits and an overall parity bit are appended to the 12 bit data, resulting in a 24 bit encoded chunk.

$$G = x^{11} + x^{10} + x^6 + x^5 + x^4 + x^2 + 1 \quad (3.1)$$

The output of the Golay encoder is shown in the table below.

Field	Data	Check bits	Parity
Position	23..12	11..1	0
Length	12	11	1

Four of these 24-bit blocks are used to reconstruct the LSF.

3.4.4 Convolutional encoder

The convolutional code shall encode the input bit sequence after appending 4 tail bits at the end of the sequence. Rate of the coder is $R=1/2$ with constraint length $K=5$ [NXDN]. The encoder diagram and generating polynomials are shown below

$$G_1(D) = 1 + D^3 + D^4 \quad (3.2)$$

$$G_2(D) = 1 + D + D^2 + D^4 \quad (3.3)$$

The output from the encoder must be read alternately.




Fig. 3.4: Convolutional coder diagram

3.4.5 Code puncturing

Removing some of the bits from the convolutional coder's output is called code puncturing. The nominal coding rate of the encoder used in M17 is $1/2$. This means the encoder outputs two bits for every bit of the input data stream. To get other (higher) coding rates, a puncturing scheme has to be used.

Two different puncturing schemes are used in M17 stream mode:

1. P_1 leaving 46 from 61 encoded bits
2. P_2 leaving 11 from 12 encoded bits

Scheme P_1 is used for the initial LICH link setup info, taking 488 bits of encoded data and selecting 368 bits. The $\gcd(368, 488)$ is 8 which, when used to divide, leaves 46 and 61. A full puncture pattern requires the output be divisible by the number of encoding polynomials. For this case the full puncture matrix should have 122 entries with 92 of them being 1.

Scheme P_2 is for frames (excluding LICH chunks, which are coded differently). This takes 296 encoded bits and selects 272 of them. Every 12th bit is being punctured out, leaving 272 bits. The full matrix shall have 12 entries with 11 being 1.

The matrix P_1 can be represented more concisely by duplicating a smaller matrix with a *flattening*.

$$S = \begin{bmatrix} a & \vec{r}_1 & c \\ b & \vec{r}_2 & X \end{bmatrix} \quad (3.4)$$

$$S_{full} = \begin{bmatrix} a & \vec{r}_1 & c & b & \vec{r}_2 \\ b & \vec{r}_2 & a & \vec{r}_1 & c \end{bmatrix} \quad (3.5)$$

The puncturing schemes are defined by their partial puncturing matrices:

$$P_1 = \begin{bmatrix} 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 & 1 & 1 \end{bmatrix} \quad (3.6)$$

$$P_2 = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 0 \end{bmatrix} \quad (3.7)$$

The complete linearized representations are:

Listing 3.1: linearized puncture patterns

```
P1 = [1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0,
1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1,
0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1,
0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1,
1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1]
```

```
P2 = [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0]
```

3.4.6 Interleaving

For interleaving a Quadratic Permutation Polynomial (QPP) is used. The polynomial $\pi(x) = (45x + 92x^2) \bmod 368$ is used for a 368 bit interleaving pattern [QPP]. See appendix Table 3.1 for pattern.

3.4.7 Data decorrelator

To avoid transmitting long sequences of constant symbols (e.g. 010101...), a simple algorithm is used. All 46 bytes of type 4 bits shall be XORed with a pseudorandom, predefined stream. The same algorithm has to be used for incoming bits at the receiver to get the original data stream. See Table 2.1 for sequence.

Data Link Layer

The Data Link layer is split into two modes:

- **Packet mode** Data are sent in small bursts, on the order of 100s to 1000s of bytes at a time, after which the physical layer stops sending data. e.g. messages, beacons, etc.
- **Stream mode** Data are sent in a continuous stream for an indefinite amount of time, with no break in physical layer output, until the stream ends. e.g. voice data, bulk data transfers, etc.

When the physical layer is idle (no RF being transmitted or received), the data link defaults to packet mode.

As is the convention with other networking protocols, all values are encoded in big endian byte order.

4.1 Stream Mode

In Stream Mode, an *indefinite* amount of payload data is sent continuously without breaks in the physical layer. The *stream* is broken up into parts, called *frames* to not confuse them with *packets* sent in packet mode. Frames contain payload data interleaved with frame signalling (similar to packets). Frame signalling is contained within the **Link Information Channel (LICH)**.

4.1.1 Sync Burst

All frames are preceded by a 16-bit *synchronization burst*.

- Link setup frames shall be preceded with 0x55F7.
- Stream frames shall be preceded with 0xFF5D.
- Packet frames shall be preceded with 0x75FF.

All syncwords are type 4 bits.

These sync words are based on [Barker codes](#). The sequence 0xDF55 (symbols -3 +3 -3 -3 +3 +3 +3 +3) is reserved.

4.1.2 Link setup frame

First frame of the transmission contains full LSF data. It's called the **Link Setup Frame (LSF)**, and is not part of any superframes.

Table 4.1: Link setup frame fields

DST	48 bits	Destination address - Encoded callsign or a special number (eg. a group)
SRC	48 bits	Source address - Encoded callsign of the originator or a special number (eg. a group)
TYPE	16 bits	Information about the incoming data stream
META	112 bits	Metadata field, suitable for cryptographic metadata like IVs or single-use numbers, or non-crypto metadata like the sender's GNSS position.
CRC	16 bits	CRC for the link setup data
TAIL	4 bits	Flushing bits for the convolutional encoder that do not carry any information. Only included for RF frames, not included for IP purposes.

Table 4.2: Bitfields of type field

Bits	Meaning
0	Packet/stream indicator, 0=packet, 1=stream
1..2	Data type indicator, 01 ₂ =data (D), 10 ₂ =voice (V), 11 ₂ =V+D, 00 ₂ =reserved
3..4	Encryption type, 00 ₂ =none, 01 ₂ =AES, 10 ₂ =scrambling, 11 ₂ =other/reserved
5..6	Encryption subtype (meaning of values depends on encryption type)
7..10	Channel Access Number (CAN)
11..15	Reserved (don't care)

The fields in Table 3 (except TAIL) form initial LSF. It contains all information needed to establish M17 link. Later in the transmission, the initial LSF is divided into 6 “chunks” and transmitted beside the payload data. This allows late-joiners to reconstruct the LICH after collecting all the pieces, and start decoding the stream even though they missed the beginning of the transmission. The process of collecting full LSF takes 6 frames or 6*40 ms = 240 ms. Four TAIL bits are needed for the convolutional coder to go back to state 0, so the ending trellis position is also known.

Voice coder rate is inferred from TYPE field, bits 1 and 2.

Table 4.3: Voice coder rates for different data type indicators

Data type indicator	Voice coder rate
00 ₂	none/reserved
01 ₂	no voice
10 ₂	3200 bps
11 ₂	1600 bps

4.1.3 Subsequent frames

Table 4.4: Fields for frames other than the link setup frame

LICH	48 bits	LSF chunk, one of 6
FN	16 bits	Frame number, starts from 0 and increments every frame to a max of 0x7fff where it will then wrap back to 0. High bit set indicates this frame is the last of the stream.
PAY-LOAD	128 bits	Payload/data, can contain arbitrary data
TAIL	4 bits	Flushing bits for the convolutional encoder that don't carry any information

The most significant bit in the FN counter is used for transmission end signalling. When transmitting the last frame, it shall be set to 1 (one), and 0 (zero) in all other frames.

The payload is used so that earlier data in the voice stream is sent first. For mixed voice and data payloads, the voice data is stored first, then the data.

Table 4.5: LSF chunk structure

Bits	Content
0..39	40 bits of full LSF
40..42	A modulo 6 counter (LICH_CNT) for LSF re-assembly
43..47	Reserved

Table 4.6: Payload example 1

Codec2 encoded frame t + 0	Codec2 encoded frame t + 1
----------------------------	----------------------------

Table 4.7: Payload Example 2

Codec2 encoded frame t + 0	Mixed data t + 0
----------------------------	------------------

4.1.4 Superframes

Each frame contains a chunk of the LSF frame that was used to establish the stream. Frames are grouped into superframes, which is the group of 6 frames that contain everything needed to rebuild the original LSF packet, so that the user who starts listening in the middle of a stream (late-joiner) is eventually able to reconstruct the LSF message and understand how to receive the in-progress stream.




Fig. 4.1: Stream consisting of one superframe

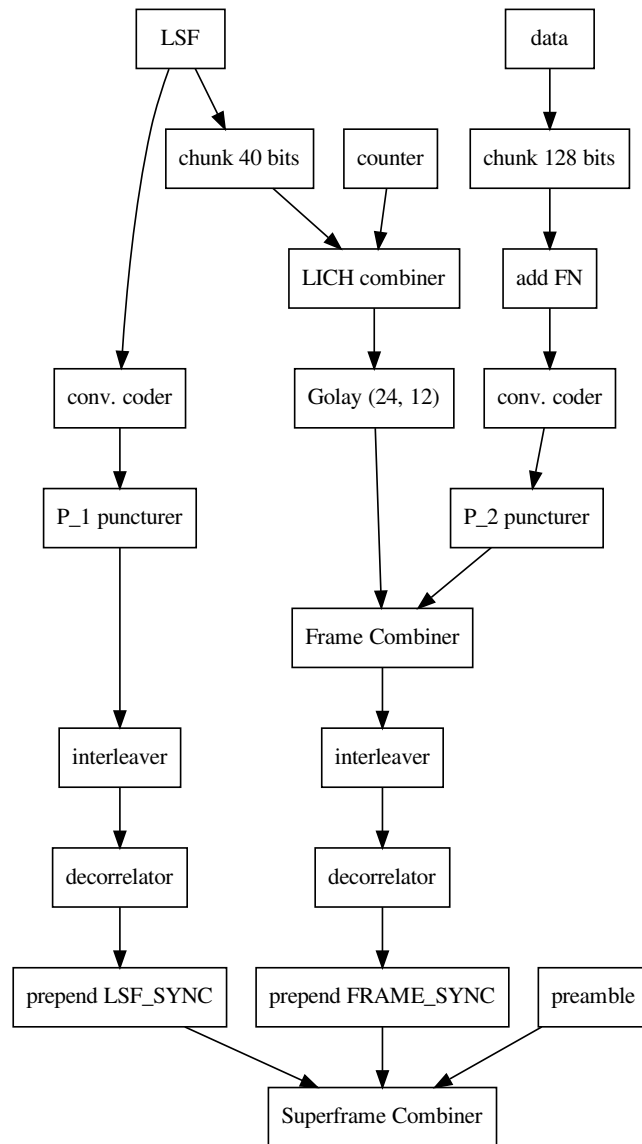


Fig. 4.2: An overview of the forward dataflow

4.1.5 CRC

M17 uses a non-standard version of 16-bit CRC with polynomial $x^{16} + x^{14} + x^{12} + x^{11} + x^8 + x^5 + x^4 + x^2 + 1$ or 0x5935 and initial value of 0xFFFF. This polynomial allows for detecting all errors up to hamming distance of 5 with payloads up to 241 bits¹, which is less than the amount of data in each frame.

As M17's native bit order is most significant bit first, neither the input nor the output of the CRC algorithm gets reflected.

The input to the CRC algorithm consists of DST, SRC (each 48 bits), 16 bits of TYPE field and 112 bits META, and then depending on whether the CRC is being computed or verified either 16 zero bits or the received CRC.

The test vectors in Table 6 are calculated by feeding the given message and then 16 zero bits to the CRC algorithm.

Table 4.8: CRC test vectors

Message	CRC output
(empty string)	0xFFFF
ASCII string "A"	0x206E
ASCII string "123456789"	0x772B
Bytes from 0x00 to 0xFF	0x1C31

4.2 Packet Mode

In *packet mode*, a finite amount of payload data (for example – text messages or application layer data) is wrapped with a packet, sent over the physical layer, and is completed when done. ~~Any acknowledgement or retransmission is done at the application layer.~~

4.2.1 Link Setup Frame

Packet mode uses the same link setup frame that has been defined for stream mode above. The packet/stream indicator is set to 0 in the type field.

Table 4.9: Bitfields of type field

Bits	Meaning
0	Packet/stream indicator, 0=packet, 1=stream
1..2	Data type indicator, 01 ₂ =data (D), 10 ₂ =voice (V), 11 ₂ =V+D, 00 ₂ =reserved
3..4	Encryption type, 00 ₂ =none, 01 ₂ =AES, 10 ₂ =scrambling, 11 ₂ =other/reserved
5..6	Encryption subtype (meaning of values depends on encryption type)
7..10	Channel Access Number (CAN)
11..15	Reserved (don't care)

Raw packet frames have no packet type metadata associated with them. Encapsulated packet format is discussed in *Packet Superframes* in the Application Layer section. This provides data type information and is the preferred format for use on M17.

When encryption type is 00₂, meaning no encryption, the encryption subtype bits are used to indicate the contents of the META field in the LSF. Since that space would otherwise go be unused, we can store small bits of data in that field such as free text or the sender's GNSS position.

Encryption type and subtype bits, including the plaintext data formats when not using encryption, are described in more detail in the Application Layer section of this document.

¹ <https://users.ece.cmu.edu/~koopman/crc/> has this listed as 0xAC9A, which is the reversed reciprocal notation

Currently the contents of the source and destination fields are arbitrary as no behavior is defined which depends on the content of these fields. The only requirement is that the content is base-40 encoded.

4.2.2 Packet Format

M17 packet mode can transmit up to 798 bytes of payload data. It achieves a base throughput of 5kbps, and a net throughput of about 4.7kbps for the largest data payload, and over 3kbps for 100-byte payloads. (Net throughput takes into account preamble and link setup overhead.)

The packet superframe consists of 798 payload data bytes and a 2-byte CCITT CRC-16 checksum.

Table 4.10: Byte fields of packet superframe

Bytes	Meaning
1-798	Packet payload
2	CCITT CRC-16

Packet data is split into frames of 368 type 4 bits preceded by a packet-specific 16-bit sync word (0xFF5D). This is the same size frame used by stream mode.

The packet frame starts with a 210 byte frame of type 1 data. It is noteworthy that it does not terminate on a byte boundary.

The frame has 200 bits (25 bytes) of payload data, 6 bits of frame metadata, and 4 bits to flush the convolutional coder.

Table 4.11: Bit fields of packet frame

Bits	Meaning
0-199	Packet payload
1	EOF indicator
5	Frame/byte count
4	Flush bits for convolutional coder

The metadata field contains a 1 bit **end of frame (EOF)** indicator, and a 5-bit frame/byte counter.

The **EOF** bit is 1 only on the last frame. The **counter** field is used to indicate the frame number when **EOF** is 0, and the number of bytes in the last frame when **EOF** is 1. This encodes the exact packet size, up to 800 bytes, in a 6-bit field.

Table 4.12: Metadata field with EOF = 0

Bits	Meaning
0	Set to 0, Not end of frame
1-5	Frame number, 0..31

Table 4.13: Metadata field with EOF = 1

Bits	Meaning
0	Set to 1, End of frame
1-5	Number of bytes in frame, 1..25

Note that it is non-conforming to send a last frame with a length of 0 bytes.

4.2.3 Convolutional Coding

The entire frame is convolutionally coded, giving 420 bits of type 2 data. It is then punctured using a 7/8 puncture matrix (1,1,1,1,1,1,0) to give 368 type 3 bits. These are then interleaved and decorrelated to give 368 type 4 bits.

Table 4.14: Packet frame

Bits	Meaning
16 bits	Sync word 0xFF5D
368 bits	Payload

4.2.4 Carrier-sense Multiple Access

When sending packets, the sender is responsible for ensuring the channel is clear before transmitting. CSMA is used to minimize collisions on a shared network. Specifically, P-persistent access is used. Each time slot is 40ms (one packet length) and the probability SHOULD default to 25%. In terms of the values used by the KISS protocol, these equate to a slot time of 4 and a P-persistence value of 63.

The benefit of this method is that it imposes no penalty on uncontested networks.

PARTS 1 AND 2 REMOVED – will add this later.

5.1 Packet Superframes

Packet superframes are composed of a 1..n byte data type specifier, 0..797 bytes of payload data. The data type specifier is encoded in the same way as UTF-8. It provides efficient coding of common data types. And it can be extended to include a very large number of distinct packet data type codes.

The data type specifier can also be used as a protocol specifier. For example, the following protocol identifiers are reserved in the M17 packet spec:

Table 5.1: Reserved Protocols

Identifier	Protocol
0x00	RAW
0x01	AX.25
0x02	APRS
0x03	6LoWPAN
0x04	IPv4
0x05	SMS
0x06	WinLink

The data type specifier is used to compute the CRC, along with the payload.

5.2 Encryption Types

Encryption is optional. The use of it may be restricted within some radio services and countries, and should only be used if legally permissible.

5.2.1 Null Encryption

Encryption type = 00_2

No encryption is performed, payload is sent in clear text.

The “Encryption SubType” bits in the Stream Type field then indicate what data is stored in the 112 bits of the LSF META field.

Encryption SubType bits	LSF META data contents
00_2	UTF-8 Text
01_2	GNSS Position Data
10_2	Reserved
11_2	Reserved

All LSF META data must be stored in big endian byte order, as throughout the rest of this specification.

GNSS Position Data stores the 112 bit META field as follows:

Size, in bits	Format	Contents
32	32-bit fixed point degrees and decimal minutes (TBD)	Latitude
32	32-bit fixed point degrees and decimal minutes (TBD)	Longitude
16	unsigned integer	Altitude, in feet MSL. Stored +1500, so a stored value of 0 represents -1500 MSL. Subtract 1500 feet when parsing.
10	unsigned integer	Course in degrees true North
10	unsigned integer	Speed in miles per hour.
12	Reserved values	Transmitter/Object description field

5.2.2 Scrambler

Encryption type = 01_2

Scrambling is an encryption by bit inversion using a bitwise exclusive-or (XOR) operation between bit sequence of data and pseudorandom bit sequence.

Encrypting bitstream is generated using a Fibonacci-topology Linear-Feedback Shift Register (LFSR). Three different LFSR sizes are available: 8, 16 and 24-bit. Each shift register has an associated polynomial. The polynomials are listed in Table 7. The LFSR is initialised with a seed value of the same length as the shift register. Seed value acts as an encryption key for the scrambler algorithm. Figures 5 to 8 show block diagrams of the algorithm

Table 5.2: LFSR scrambler polynomials

Encryption subtype	LFSR polynomial	Seed length	Sequence period
00 ₂	$x^8 + x^6 + x^5 + x^4 + 1$	8 bits	255
01 ₂	$x^{16} + x^{15} + x^{13} + x^4 + 1$	16 bits	65,535
10 ₂	$x^{24} + x^{23} + x^{22} + x^{17} + 1$	24 bits	16,777,215




Fig. 5.1: 8-bit LFSR taps




Fig. 5.2: 16-bit LFSR taps




Fig. 5.3: 24-bit LFSR taps

5.2.3 Advanced Encryption Standard (AES)

Encryption type = 10₂

This method uses AES block cipher in counter (CTR) mode, with a 96-bit nonce that should never be used for more than one separate stream and a 32 bit CTR.

The 96-bit AES nonce value is extracted from the 96 most significant bits of the META field, and the remaining 16 bits of the META field form the highest 16 bits of the 32 bit counter. The FN (Frame Number) field value is then used to fill out the lower 16 bits of the counter, and always starts from 0 (zero) in a new voice stream.

The 16 bit frame number and 40 ms frames can provide for over 20 minutes of streaming without rolling over the counter¹.

The random part of the nonce value should be generated with a hardware random number generator or any other method of generating non-repeating values.

To combat replay attacks, a 32-bit timestamp shall be embedded into the cryptographic nonce field. The field structure of the 96 bit nonce is shown in Table 9. Timestamp is 32 LSB portion of the number of seconds that elapsed since the beginning of 1970-01-01, 00:00:00 UTC, minus leap seconds (a.k.a. “unix time”).

¹ The effective capacity of the counter is 15 bits, as the MSB is used for transmission end signalling. At 40ms per frame, or 25 frames per second, and 2**15 frames, we get 2**15 frames / 25 frames per second = 1310 seconds, or 21 minutes and some change.

Table 5.3: 96 bit nonce field structure

TIMESTAMP	RANDOM DATA	CTR_HIGH
32	64	16

CTR_HIGH field initializes the highest 16 bits of the CTR, with the rest of the counter being equal to the FN counter.

Warning: In CTR mode, AES encryption is malleable [CTR] [CRYPTO]. That is, an attacker can change the contents of the encrypted message without decrypting it. This means that recipients of AES-encrypted data must not trust that the data is authentic. Users who require that received messages are proven to be exactly as-sent by the sender should add application-layer authentication, such as HMAC. In the future, use of a different mode, such as Galois/Counter Mode, could alleviate this issue [CRYPTO].

APPENDIX A

Address Encoding

M17 uses 48 bits (6 bytes) long addresses. Callsigns (and other addresses) are encoded into these 6 bytes in the following ways:

- An address of 0 is invalid.
- Address values between 1 and 26214399999999 (which is $40^9 - 1$), up to 9 characters of text are encoded using base40, described below.
- Address values between 262144000000000 (40^9) and 281474976710654 ($2^{48} - 2$) are invalid
- An address of 0xFFFFFFFF is a broadcast. All stations should receive and listen to this message.

Table 1.1: Address scheme

Address Range	Category	Number of addresses	Remarks
0x000000000000	RESERVED	1	For future use
0x000000000001-0xee6b27ffff	Unit ID	26214399999999	
0xee6b28000000-0xfffffffffe	RESERVED	19330976710655	For future use
0xffffffff	Broadcast	1	Valid only for destination field

A.1 Callsign Encoding: base40

9 characters from an alphabet of 40 possible characters can be encoded into 48 bits, 6 bytes. The base40 alphabet is:

- 0: A space. Invalid characters will be replaced with this.
- 1-26: “A” through “Z”
- 27-36: “0” through “9”
- 37: “-” (hyphen)
- 38: “/” (slash)
- 39: “.” (dot)

Encoding is little endian. That is, the right most characters in the encoded string are the most significant bits in the resulting encoding.

A.1.1 Example code: encode_base40()

```
uint64_t encode_callsign_base40(const char *callsign) {
    uint64_t encoded = 0;
    for (const char *p = (callsign + strlen(callsign) - 1); p >= callsign; p-- ) {
        encoded *= 40;
        // If speed is more important than code space,
        // you can replace this with a lookup into a 256 byte array.
        if (*p >= 'A' && *p <= 'Z') // 1-26
            encoded += *p - 'A' + 1;
        else if (*p >= '0' && *p <= '9') // 27-36
            encoded += *p - '0' + 27;
        else if (*p == '-') // 37
            encoded += 37;
        // These are just place holders. If other characters make more sense,
        // change these. Be sure to change them in the decode array below too.
        else if (*p == '/') // 38
            encoded += 38;
        else if (*p == '.') // 39
            encoded += 39;
        else
            // Invalid character or space, represented by 0, decoded as a space.
            //encoded += 0;
    }
    return encoded;
}
```

A.1.2 Example code: decode_base40()

```
char *decode_callsign_base40(uint64_t encoded, char *callsign) {
    if (encoded >= 262144000000000) { // 40^9
        *callsign = 0;
        return callsign;
    }
    char *p = callsign;
    for (; encoded > 0; p++) {
        *p = " ABCDEFGHIJKLMNOPQRSTUVWXYZ0123456789-/. "[encoded % 40];
        encoded /= 40;
    }
    *p = 0;

    return callsign;
}
```

A.1.3 Why base40?

The longest commonly assigned callsign from the FCC is 6 characters. The minimum alphabet of A-Z, 0-9, and a “done” character mean the most compact encoding of an American callsign could be: $\log_2(37^6) = 31.26$ bits, or 4 bytes.

Some countries use longer callsigns, and the US sometimes issues longer special event callsigns. Also, we want to extend our callsigns (see below). So we want more than 6 characters. How many bits do we need to represent more characters:

Table 1.2: bits per characters

characters	bits	bytes
7	$\log_2(37^7) = 36.47$	5
8	$\log_2(37^8) = 41.67$	6
9	$\log_2(37^9) = 46.89$	6
10	$\log_2(37^{10}) = 52.09$	7

Of these, 9 characters into 6 bytes seems the sweet spot. Given 9 characters, how large can we make the alphabet without using more than 6 bytes?

Table 1.3: alphabet size vs bytes

alphabet size	bits	bytes
37	$\log_2(37^9) = 46.89$	6
38	$\log_2(38^9) = 47.23$	6
39	$\log_2(39^9) = 47.57$	6
40	$\log_2(40^9) = 47.90$	6
41	$\log_2(41^9) = 48.22$	7

Given this, 9 characters from an alphabet of 40 possible characters, makes maximal use of 6 bytes.

A.2 Callsign Formats

Government issued callsigns should be able to encode directly with no changes.

A.2.1 Multiple Stations

To allow for multiple stations by the same operator, we borrow the use of the ‘-’ character from AX.25 and the SSID field. A callsign such as “AB1CD-1” is considered a different station than “AB1CD-2” or even “AB1CD”, but it is understood that these all belong to the same operator, “AB1CD”

A.2.2 Temporary Modifiers

Similarly, suffixes are often added to callsign to indicate temporary changes of status, such as “AB1CD/M” for a mobile station, or “AB1CD/AE” to signify that I have Amateur Extra operating privileges even though the FCC database may not yet be updated. So the ‘/’ is included in the base40 alphabet. The difference between ‘-’ and ‘/’ is that ‘-’ are considered different stations, but ‘/’ are NOT. They are considered to be a temporary modification to the same station.

A.2.3 Interoperability

It may be desirable to bridge information between M17 and other networks. The 9 character base40 encoding allows for this:

DMR

DMR unfortunately doesn’t have a guaranteed single name space. Individual IDs are reasonably well recognized to be managed by <https://www.radioid.net/database/search#!> but Talk Groups are much less well managed. Talk Group XYZ on Brandmeister may be (and often is) different than Talk Group XYZ on a private cBridge system.

- DMR IDs are encoded as: D<number> eg: D3106728 for KR6ZY
- DMR Talk Groups are encoded by their network. Currently, the following networks are defined:
- Brandmeister: BM<number> eg: BM31075
- DMRPlus: DP<number> eg: DP262
- More networks to be defined here.

D-Star

D-Star reflectors have well defined names: REFxxxY which are encoded directly into base40.

APPENDIX B

Decorrelator sequence

Table 2.1: Decorrelator scrambling sequence

Seq. number	Value	Seq. number	Value
00	0xD6	23	0x6E
01	0xB5	24	0x68
02	0xE2	25	0x2F
03	0x30	26	0x35
04	0x82	27	0xDA
05	0xFF	28	0x14
06	0x84	29	0xEA
07	0x62	30	0xCD
08	0xBA	31	0x76
09	0x4E	32	0x19
10	0x96	33	0x8D
11	0x90	34	0xD5
12	0xD8	35	0x80
13	0x98	36	0xD1
14	0xDD	37	0x33
15	0x5D	38	0x87
16	0x0C	39	0x13
17	0xC8	40	0x57
18	0x52	41	0x18
19	0x43	42	0x2D
20	0x91	43	0x29
21	0x1D	44	0x78
22	0xF8	45	0xC3

Interleaving

Table 3.1: Interleaving table

input index	output ind	input index	output ind	input index	output ind	input index	output ind
0	0	92	92	184	184	276	276
1	137	93	229	185	321	277	45
2	90	94	182	186	274	278	366
3	227	95	319	187	43	279	135
4	180	96	272	188	364	280	88
5	317	97	41	189	133	281	225
6	270	98	362	190	86	282	178
7	39	99	131	191	223	283	315
8	360	100	84	192	176	284	268
9	129	101	221	193	313	285	37
10	82	102	174	194	266	286	358
11	219	103	311	195	35	287	127
12	172	104	264	196	356	288	80
13	309	105	33	197	125	289	217
14	262	106	354	198	78	290	170
15	31	107	123	199	215	291	307
16	352	108	76	200	168	292	260
17	121	109	213	201	305	293	29
18	74	110	166	202	258	294	350
19	211	111	303	203	27	295	119
20	164	112	256	204	348	296	72
21	301	113	25	205	117	297	209
22	254	114	346	206	70	298	162
23	23	115	115	207	207	299	299
24	344	116	68	208	160	300	252
25	113	117	205	209	297	301	21
26	66	118	158	210	250	302	342

Continued on next page

Table 3.1 – continued from previous page

input index	output ind	input index	output ind	input index	output ind	input index	output ind
27	203	119	295	211	19	303	111
28	156	120	248	212	340	304	64
29	293	121	17	213	109	305	201
30	246	122	338	214	62	306	154
31	15	123	107	215	199	307	291
32	336	124	60	216	152	308	244
33	105	125	197	217	289	309	13
34	58	126	150	218	242	310	334
35	195	127	287	219	11	311	103
36	148	128	240	220	332	312	56
37	285	129	9	221	101	313	193
38	238	130	330	222	54	314	146
39	7	131	99	223	191	315	283
40	328	132	52	224	144	316	236
41	97	133	189	225	281	317	5
42	50	134	142	226	234	318	326
43	187	135	279	227	3	319	95
44	140	136	232	228	324	320	48
45	277	137	1	229	93	321	185
46	230	138	322	230	46	322	138
47	367	139	91	231	183	323	275
48	320	140	44	232	136	324	228
49	89	141	181	233	273	325	365
50	42	142	134	234	226	326	318
51	179	143	271	235	363	327	87
52	132	144	224	236	316	328	40
53	269	145	361	237	85	329	177
54	222	146	314	238	38	330	130
55	359	147	83	239	175	331	267
56	312	148	36	240	128	332	220
57	81	149	173	241	265	333	357
58	34	150	126	242	218	334	310
59	171	151	263	243	355	335	79
60	124	152	216	244	308	336	32
61	261	153	353	245	77	337	169
62	214	154	306	246	30	338	122
63	351	155	75	247	167	339	259
64	304	156	28	248	120	340	212
65	73	157	165	249	257	341	349
66	26	158	118	250	210	342	302
67	163	159	255	251	347	343	71
68	116	160	208	252	300	344	24
69	253	161	345	253	69	345	161
70	206	162	298	254	22	346	114
71	343	163	67	255	159	347	251
72	296	164	20	256	112	348	204
73	65	165	157	257	249	349	341
74	18	166	110	258	202	350	294
75	155	167	247	259	339	351	63

Continued on next page

Table 3.1 – continued from previous page

input index	output ind	input index	output ind	input index	output ind	input index	output ind
76	108	168	200	260	292	352	16
77	245	169	337	261	61	353	153
78	198	170	290	262	14	354	106
79	335	171	59	263	151	355	243
80	288	172	12	264	104	356	196
81	57	173	149	265	241	357	333
82	10	174	102	266	194	358	286
83	147	175	239	267	331	359	55
84	100	176	192	268	284	360	8
85	237	177	329	269	53	361	145
86	190	178	282	270	6	362	98
87	327	179	51	271	143	363	235
88	280	180	4	272	96	364	188
89	49	181	141	273	233	365	325
90	2	182	94	274	186	366	278
91	139	183	231	275	323	367	47

M17 Internet Protocol (IP) Networking

Digital modes are commonly networked together through linked repeaters using IP networking.

For commercial protocols like DMR, this is meant for linking metropolitan and state networks together and allows for easy interoperability between radio users. Amateur Radio uses this capability for creating global communications networks for all imaginable purposes, and makes ‘working the world’ with an HT possible.

M17 is designed with this use in mind, and has native IP framing to support it.

In competing radio protocols, a repeater or some other RF to IP bridge is required for linking, leading to the use of hotspots (tiny simplex RF bridges).

The TR-9 and other M17 radios may support IP networking directly, such as through the ubiquitous ESP8266 chip or similar. This allows them to skip the RF link that current hotspot systems require, finally bringing to fruition the “Amateur digital radio is just VoIP” dystopian future we were all warned about.

D.1 Standard IP Framing

M17 over IP is big endian, consistent with other IP protocols. We have standardized on UDP port 17000, this port is recommended but not required. Later specifications may require this port.

Table 4.1: Internet frame fields

MAGIC	32 bits	Magic bytes 0x4d313720 (“M17 “)
StreamID (SID)	16 bits	Random bits, changed for each PTT or stream, but consistent from frame to frame within a stream
LICH	240 bits	The meaningful contents of a LICH frame (dst, src, streamtype, META field, CRC16) as defined earlier.
FN	16 bits	Frame number (exactly as would be transmitted as an RF stream frame, including the last frame indicator at (FN & 0x8000))
Payload	128 bits	Payload (exactly as would be transmitted in an RF stream frame)
CRC16	16 bits	CRC for the entire packet, as defined earlier (TODO: specific link)

The CRC checksum must be recomputed after modification or re-assembly of the packet, such as when translating from RF to IP framing.

D.2 Control Packets

Reflectors use a few different types of control frames, identified by their magic:

- *CONN* - Connect to a reflector
- *ACKN* - acknowledge connection
- *NACK* - deny connection
- *PING* - keepalive for the connection from the reflector to the client
- *PONG* - keepalive response from the client to the reflector
- *DISC* - Disconnect (client->reflector or reflector->client)

D.2.1 CONN

Table 4.2: Bytes of a CONN packet

Bytes	Purpose
0-3	Magic - ASCII “CONN”
4-9	6-byte ‘From’ callsign including module in last character (e.g. “A1BCD D”) encoded as per <i>Address Encoding</i>
10	Module to connect to - single ASCII byte A-Z

A client sends this to a reflector to initiate a connection. The reflector replies with ACKN on successful linking, or NACK on failure.

D.2.2 ACKN

Table 4.3: Bytes of ACKN packet

Bytes	Purpose
0-3	Magic - ASCII “ACKN”

D.2.3 NACK

Table 4.4: Bytes of NACK packet

Bytes	Purpose
0-3	Magic - ASCII “NACK”

D.2.4 PING

Table 4.5: Bytes of PING packet

Bytes	Purpose
0-3	Magic - ASCII “PING”
4-9	6-byte ‘From’ callsign including module in last character (e.g. “A1BCD D”) encoded as per <i>Address Encoding</i>

Upon receiving a PING from a reflector, the client replies with a PONG

D.2.5 PONG

Table 4.6: Bytes of PONG packet

Bytes	Purpose
0-3	Magic - ASCII “PONG”
4-9	6-byte ‘From’ callsign including module in last character (e.g. “A1BCD D”) encoded as per <i>Address Encoding</i>

D.2.6 DISC

Table 4.7: Bytes of DISC packet

Bytes	Purpose
0-3	Magic - ASCII “DISC”
4-9	6-byte ‘From’ callsign including module in last character (e.g. “A1BCD D”) encoded as per <i>Address Encoding</i>

Sent by either end to force a disconnection. Acknowledged with 4-byte packet “DISC” (without the callsign field)

APPENDIX E

KISS Protocol

The purpose of this appendix is to document conventions for adapting KISS TNCs to M17 packet and streaming modes. M17 is a more complex protocol, both at the baseband level and at the data link layer than is typical for HDLC-based protocols commonly used on KISS TNCs. However, it is well suited for modern packet data links, and can even be used to stream digital audio between a host and a radio.

This appendix assumes the reader is familiar with the streaming and packet modes defined in the M17 spec, and with KISS TNCs and the KISS protocol.

In all cases, the TNC expects to get the data payload to be sent and is responsible for frame construction, FEC encoding, puncturing, interleaving and decorrelation. It is also responsible for baseband modulation.

For streaming modes, all voice encoding (Codec2) is done on the host and not on the TNC. The host is also responsible for constructing the LICH.

E.1 References

- <http://www.ax25.net/kiss.aspx>
- <https://packet-radio.net/wp-content/uploads/2017/04/multi-kiss.pdf>
- https://en.wikipedia.org/wiki/OSI_model

E.2 Glossary

TNC Terminal node controller – a baseband network interface device to allow host computers to send data over a radio network, similar to a modem. It connects a computer to a radio and handles the baseband portion of the physical layer and the data link layer of network protocol stack.

KISS Short for “Keep it simple, stupid”. A simplified TNC protocol designed to move everything except for the physical layer and the data link layer out of the TNC. Early TNCs could include everything up through the application layer of the OSI network model.

SLIP *Serial Line Internet Protocol* – the base protocol used by the KISS protocol, extended by adding a single *type indicator* byte at the start of a frame.

type indicator A one byte code at the beginning of a KISS frame which indicates the TNC *port* and KISS *command*.

port A logical port on a TNC. This allowed a single TNC to connect to multiple radios. Its specific use is loosely defined in the KISS spec. The high nibble of the KISS *type indicator*. Port 0xF is reserved.

command A KISS command. This tells the TNC or host how to interpret the KISS frame contents. The low nibble of the KISS *type indicator*. Command 0xF is reserved.

CSMA *Carrier-sense multiple access* – a protocol used by network devices to minimize collisions on a shared communications channel.

HDLC *High-Level Data Link Control* – a data link layer framing protocol used in many AX.25 packet radio networks. Many existing protocol documents, including KISS, reference HDLC because of its ubiquity when the protocols were invented. However, HDLC is not a requirement for higher level protocols like KISS which are agnostic to the framing used at the data link layer.

EOS End of stream – an indicator bit in the frame number field of a stream data frame.

LICH Link information channel – a secondary data channel in the stream data frame containing supplemental information, including a copy of the link setup frame.

E.3 M17 Protocols

This specification defines KISS TNC modes for M17 packet and streaming modes, allowing the KISS protocol to be used to send and receive M17 packet and voice data. Both are bidirectional. There are two packet modes defined. This is done to provide complete access to the M17 protocol while maintaining the greatest degree of backwards compatibility with existing packet applications.

These protocols map to specific KISS *port*. The host tells the TNC what type of data to transmit based on the port used in host to TNC transfers. And the TNC tells the host what data it has received by the port set on TNC to host transfers.

This document outlines first the two packet protocols, followed by the streaming protocol.

E.4 KISS Basics

E.4.1 TX Delay

If a KISS **TX delay** T_d greater than 0 is specified, the transmitter is keyed for $T_d * 10ms$ with only a DC signal present.

The T_d value should be adjusted to the minimum required by the transmitter in order to transmit the full preamble reliably.

Only a single 40ms preamble frame is ever sent.

Note: A TX delay may be necessary because many radios require some time between when PTT is engaged and the transmitter can begin transmitting a modulated signal.

E.5 Packet Protocols

In order to provide backward compatibility with the widest range of existing ham radio software, and to make use of features in the the M17 protocol itself, we will define two distinct packet interfaces *BASIC* and *FULL*.

The KISS protocol allows us to target specific modems using the port identifier in the control byte.

We first define basic packet mode as this is initially likely to be the most commonly used mode over KISS.

E.5.1 M17 Basic Packet Mode

Basic packet mode uses only the standard KISS protocol on **TNC port 0**. This is the default port for all TNCs. Packets are sent using command 0. Again, this is normal behavior for KISS client applications.

Sending Data

In basic mode, the TNC only expects to receive packets from the host, as it would for any other mode supported AFSK, G3RUH, etc.

If the TNC is configured for half-duplex, the TNC will do P-persistence CSMA using a 40ms slot time and obey the P value set via the KISS interface. CSMA is disabled in full-duplex mode.

The **TX Tail** value is deprecated and is ignored.

The TNC sends the preamble burst.

The TNC is responsible for constructing the link setup frame, identifying the content as a raw mode packet. The source field is an encoded TNC identifier, similar to the APRS TOCALL, but it can be an arbitrary text string up to 9 characters in length. The destination is set to the broadcast address.

In basic packet mode, it is expected that the sender callsign is embedded within the packet payload.

The TNC sends the link setup frame.

The TNC then computes the CRC for the full packet, splits the packet into data frames encode and modulate each frame back-to-back until the packet is completely transmitted.

If there is another packet to be sent, the preamble can be skipped and the TNC will construct the next link setup frame (it can re-use the same link setup frame as it does not change) and send the next set of packet frames.

Limitations

The KISS specification defines no limitation to the packet size allowed. Nor does it specify any means of returning error conditions back to the host. M17 packet protocol limits the raw packet payload size to 798 bytes. The TNC must drop any packets larger than this.

Receiving Data

When receiving M17 data, the TNC must receive and parse the link setup frame and verify that the following frames contain raw packet data.

The TNC is responsible for decoding each packet, assembling the packet from the sequence of frames received, and verifying the packet checksum. If the checksum is valid, the TNC transfers the packet, excluding the CRC to the host using **KISS port 0**.

E.5.2 M17 Full Packet Mode

The purpose of full packet mode is to provide access to the entire M17 packet protocol to the host. This allows the host to set the source and destination fields, filter received packets based on the content these fields, enable encryption, and send and receive type-coded frames.

Use M17 full packet mode by sending to **KISS port 1**. In this mode the host is responsible for sending both the link setup frame and the packet data. It does this by prepending the 30-byte link setup frame to the packet data, sending this to the TNC in a single KISS frame. The TNC uses the first 30 bytes as the link setup frame verbatim, then splits the remaining data into M17 packet frames.

As with basic mode, the TNC uses the **Duplex** setting to enable/disable CSMA, and uses the **P value** for CSMA, with a fixed slot time of “4” (40 ms).

Receiving Data

For TNC to host transfers, the same occurs. The TNC combines the link setup frame with the packet frame and sends both in one KISS frame to the host using **KISS port 1**.

E.6 Stream Protocol

The streaming protocol is fairly trivial to describe. It is used by sending first a link setup frame followed by a stream of 26-byte data frames to **KISS port 2**.

E.6.1 Stream Format

Table 5.1: M17 KISS Stream Protocol

Frame Size	Contents
30	Link Setup Frame
26	LICH + Payload
26	LICH + Payload
...	...
26	LICH + Payload with EOS bit set.

The host must not send any frame to any other KISS port while a stream is active (a frame with the EOS bit has not been sent).

It is a protocol violation to send anything other than a link setup frame with the stream mode bit set in the first field as the first frame in a stream transfer to KISS port 2. Any such frame is ignored.

It is a protocol violation to send anything to any other KISS port while a stream is active. If that happens the stream is terminated and the packet that caused the protocol violation is dropped.

E.6.2 Data Frames

The data frames contain a 6-byte (48-bit) LICH segment followed by a 20 byte payload segment consisting of frame number, 16-byte data payload and CRC. The TNC is responsible for parsing the frame number and detecting the end-of-stream bit to stop transmitting.

Table 5.2: KISS Stream Data Frame

Frame Size	Contents
6	LICH (48 bits)
2	Frame number and EOS flag
16	Payload
2	M17 CRC of frame number and payload

The TNC is responsible for FEC-encoding both the LICH the payload, as well as interleaving, decorrelation, and baseband modulation.

E.6.3 Timing Constraints

Streaming mode provides additional timing constraints on both host to TNC transfers and on TNC to host transfers. Payload frames must arrive every 40ms and must have a jitter below 40ms. In general, it is expected that the TNC has up to 2 frames buffered (buffering occurs while sending the preamble and link setup frames), it should be able to keep the transmit buffers filled with packet jitter of 40ms.

The TNC must stop transmitting if the transmit buffers are empty. The TNC communicates that it has stopped transmitting early (before seeing a frame with the **end of stream** indicator set) by sending an empty data frame to the host.

E.6.4 TNC to Host Transfers

TNC to host transfers are similar in that the TNC first sends the 30-byte link setup frame received to the host, followed by a stream of 26-byte data frames as described above. These are sent using **KISS port 2**.

The TNC must send the link setup frame first. This means that the TNC must be able to decode LICH segments and assemble a valid link setup frame before it sends the first data frame. The TNC will only send a link setup frame with a valid CRC to the host. After the link setup frame is sent, the TNC ignores the CRC and sends all valid frames (those received after a valid sync word) to the host. If the stream is lost before seeing an end-of-stream flag, the TNC sends a 0-byte data frame to indicate loss of signal.

The TNC must then re-acquire the signal by decoding a valid link setup frame from the LICH in order to resume sending to the host.

E.6.5 Busy Channel Lockout

The TNC implements **busy channel lockout** by enabling half-duplex mode on the TNC, and disables **busy channel lockout** by enabling full-duplex mode. When busy channel lockout occurs, the TNC keeps the link setup frame and discards all data frames until the channel is available. It then sends the preamble, link setup frame, and starts sending the data frames as they are received.

Note: BCL will be apparent to a receiver as the first frame received after the link setup frame will not start with frame number 0.

E.6.6 Limitations

Information is lost by having the TNC decode the LICH. It is not possible to communicate to the host that the LICH bytes are known to be invalid.

Should we have the TNC signal the host by dropping known invalid LICH segments? The host can tell that the LICH is missing by looking at the frame size.

E.7 Mixing Modes

An M17 KISS TNC need not keep track of state across distinct TNC ports. Packet transfers are sent one packet at a time. It is OK to send to port 0 and port 1 in subsequent transfers. It is also OK to send a packet followed immediately by a voice streams. As mentioned earlier, it is a protocol violation to send a KISS frame to any other port while a stream is active. However, a packet can be sent immediately following a voice stream (after EOS is sent).

E.7.1 Back-to-back Transfers

The TNC is expected to detect back-to-back transfers from the host, even across different KISS ports, and suppress the generation of the preamble.

For example, a packet containing APRS data sent immediately on PTT key-up should be sent immediately after the EOS frame.

Back-to-back transfers are common for packet communication where the **window size** determines the number of unacknowledged frames which may be outstanding (unacknowledged). Packet applications will frequently send back-to-back packets (up to **window size** packets) before waiting for the remote end to send ACKs for each of the packets.

E.8 Implementation Details

E.8.1 Polarity

One of the issues that must be addressed by the TNC designer, and one which the KISS protocol offers no ready solution for, is the issue of polarity.

A TNC must interface with a RF transceiver for a complete M17 physical layer implementation. RF transceivers may have different polarity for their TX and RX paths.

M17 defines that the +3 symbol is transmitted with a +2.4 kHz deviation (2.4 kHz above the carrier). **Normal polarity** in a transceiver results in a positive voltage driving the frequency higher and a lower voltage driving the frequency lower. **Reverse polarity** is the opposite. A higher voltage drives the frequency lower.

On the receive side the same issue exists. **Normal polarity** results in a positive voltage output when the received signal is above the carrier frequency. **Reverse polarity** results in a positive voltage when the frequency is below the carrier.

Just as with transmitter deviation levels and received signal levels, the polarity of the transmit and receive path must be adjustable on a 4-FSK modem. The way these adjustments are made to the TNC are not addressed by the KISS specification.

Bibliography

- [TETRA] Dunlop, John; Girma, Demessie; Irvine, James “Digital Mobile Communications and the TETRA System” Wiley 1999, ISBN: 9780471987925
- [DMR] ETSI TS 102 361-1 V2.2.1 (2013-02): “Electromagnetic compatibility and Radio spectrum Matters (ERM); Digital Mobile Radio (DMR) Systems; Part 1: DMR Air Interface (AI) protocol” https://www.etsi.org/deliver/etsi_ts/102300_102399/10236101/02.02.01_60/ts_10236101v020201p.pdf
- [ECC] Moreira, Jorge C.; Farrell, Patrick G. “Essentials of Error-Control Coding” Wiley 2006, ISBN: 9780470029206
- [NXDN] NXDN Technical Specifications, Part 1: Air Interface; Sub-part A: Common Air Interface
- [QPP] Trifina, Lucian, Daniela Tarniceriu, and Valeriu Munteanu. “Improved QPP Interleavers for LTE Standard.” ISSCS 2011 - International Symposium on Signals, Circuits and Systems (2011): n. pag. Crossref. Web. <https://arxiv.org/abs/1103.3794>
- [CTR] McGrew, David A. “Counter mode security: Analysis and recommendations.” Cisco Systems, November 2, no. 4 (2002).
- [CRYPTO] Rogaway, Phillip. “Evaluation of some blockcipher modes of operation.” Cryptography Research and Evaluation Committees (CRYPTREC) for the Government of Japan (2011).

C

command, [38](#)
CSMA, [38](#)

E

ECC, [3](#)
EOS, [38](#)

F

FEC, [3](#)
Frame, [3](#)

H

HDLC, [38](#)

K

KISS, [37](#)

L

LICH, [3](#), [38](#)
Link Information Frame, [3](#)

P

Packet, [3](#)
port, [38](#)

S

SLIP, [38](#)
Superframe, [3](#)

T

TNC, [37](#)
type indicator, [38](#)